

Multivariate Time Series Classification and Embedding

David Pham

September 16, 2018

Contents

1	Project overview	1
2	Problem Statement	2
3	Metrics	2
4	Data	2
5	Exploratory Analysis	3
6	Algorithms and Technique	3
7	Benchmark	5
8	Methodology	5
8.1	Data processing	5
8.2	Implementation	5
9	Refinement	6
9.1	Structure of the network	6
10	Results	6
10.1	Benchmark comparison	11
11	Embedding	11
12	Conclusion and Further applications	14

1 Project overview

In the financial investment industry, the speed and the quality of processing information is what determines if a business is successful or failure. In the recent period, an emphasis has been set on machine learning algorithms for investment strategies. One common application is the statement of the algorithms to perform classification of companies to put them into sector.

This project aims to create and review algorithms for classifying companies by using their share price. We will test a heuristic classifier using sample correlation and train a deep neural network to mimic the performance the sample correlation and study the hidden layer of the network. The projects uses data from Quandl and SimFin.

Time series modeling using deep learning in a financial context has been studied by Bao and Yue (2017), where share prices are predicted using a combination of wavelet transform, autoencoders and long-short term memory units. Qiu et al. (2018) studies how to create

fast inference of correlation of between time series by applying discrete Fourier transform and uses the frequency domain of the time series as embedding for the input. A dense neural network is then trained to create a model that predict correlations between embedding of the time series. Several loss function are studied and lead to different embedding. Our work differs as our methods keeps in time domain of the time series and aims to classify them more as to find the nearest neighbors and to predict them accurately and efficiently. Additional reference can be found in the aforementioned papers.

2 Problem Statement

The goal of this project is to create a method to embed financial time series into a vector space and analyze qualitatively the embedding. The main challenge to complete this projects are following.

1. Filter the universe of the data and acquire the data from different sources.
2. Transform data to be usable in machine learning algorithm.
3. Create artificial data to transform the problem into a supervised learning problem.
4. Create a base model and train an advanced classifier to create the embedding.
5. Analyze briefly the embedding.

3 Metrics

We will use the *accuracy* defined as the ratio between the true positive and negatives predictions divided by the dataset size. Or in formula

$$\text{accuracy} = \frac{\text{NUMBER OF CORRECT CLASSIFICATION}}{\text{SAMPLE SIZE}}$$

This metric is relevant because at a certain level the algorithms should be able to assign correctly the stock price of companies to the sector of the companies.

Compared to other metrics, the accuracy targets our ability to predict the sector of each companies. As our model aims to be as general as possible, there is no particular additional cost of having false negative or false positive, hence the accuracy is a suitable metric for metric.

4 Data

The below table shows the data and how to access them.

Type	Provider	Access to Datatest	Downloaded
Price	Quandl	Quandl API	2018-08-21
Company information	SimFin	S3	2018-08-21
Stocks Returns	Processed	S3	2018-09-11
Index Returns	Processed	S3	2018-09-11

In our analysis, we used the closed price of the US companies provided by Wikipedia and distributed by Quandl. As we additionally required the sectors from the company we needed to merge it with an additional databases provided by SimFin, which provides the taxonomy created by Global Industry Classification Standard (GICS). We could find the sectors of 726 companies which constituted our dataset. The prices range from 1962 to 2018.




Figure 1: Distribution of sectors in the data.

The stocks returns have a mean 7 basis points (0.0007) and median of 0, whereas the standard deviation is at 10 percent. When the returns are bounded between the 5th and 95th quantile, they range between -3.6% and 3.8% which are loosely normally distributed with a standard deviation estimated at 1.5%. The returns display some maximum values that are abnormal (like a share price multiplication by 99, which can be possible but are exceptional).

Concerning the indices returns, by construction their min and max values are set to -10% and 10%. When filtering the indices returns lying between the 5th and 95th quantile, they range between -2.1% and 2.2% and seems to follow a Gaussian distribution with mean 0 and standard deviation of 1.5%.

5 Exploratory Analysis

In Figure 5 and 5, one can observe that number of stocks per sectors and the synthetic indices created by aggregating the returns.

We observe that financial services, technology and oil and gas are the most prominent sectors. At first glance, the indices seems to be behave according to our intuition to the economic situation from the last decades. We see clear two downward pikes which coincide with the bubble dot-com and the financial crisis.

6 Algorithms and Technique

There is two classifiers. One is algorithmic and the is other one is based on a dense neural network using the right transformation of the input to incorporate non linear relationship between the weights.

The first classifier computes correlations of time-series and then assign to an input time-series the *sector* with which it has the highest correlation.

For two random variables X and Y , with sufficient stability assumption, the Pearson correlation is defined as




Figure 2: Synthetic indices of sectors according to GICS. Index are set on 100 on the 1990-01-01.

$$\rho_p(X, Y) = E[XY] - (E[X]E[Y])^2 \approx \sum_{i=1}^n x_i y_i - \left(\sum_{i=1}^n x_i \sum_{i=1}^n y_i \right)^2,$$

This first rule based method does not require any parameter fitting (which makes it appealing with some respect).

The second classifier use a deep neural network to create features and representation (or state) of the inputs to be able to separate linearly all the output sectors (or classes).

A deep neural network can be understood as a powerful mapping f_θ between the input and the output. In our particular case, f_θ can be interpreted as function composition of many simpler models whose output have been transformed by a non-linear function (called activation function). The goal of these composition is to support the network to design by himself relevant transformations of the input, called embedding or latent variables, to help it improving its predictions.

The challenge of the designer of the network is to specify the type and size of these inner simpler function and the non linear function linking them to create a model whose input transformations will generalize to out of sample data.

The model then predicts its output by using these transformations of the input spaces by creating partition of the space. A good intuition of the last step can gain on Playground Tensorflow.

Another challenge is to be able to find optimal weights $\hat{\theta}$ for f_θ . This problem is solved by running optimization algorithms trying to reduce a loss function. The choice of the loss function and the optimization algorithm are key decision for the training of the model.

We use the *categorical cross entropy* loss in our model, which basically minimize the probability of the model to pick the wrong sector. We train the neural network using *Adam*, which optimize the parameters of the network by minimizing the loss using stochastic

gradient descent algorithm which standardizes the gradient and also updates it in iteration. In detail, it means that our optimizer creates estimates of θ and update θ of f_θ in the direction such that the loss diminishes the most. The size of the update is named the learning rate. The direction, also called the gradient, is estimated on a sample of observations, the so-called batch, and *Adam* creates a weighted average the latest estimator of the gradient and the new (standardized) estimator. This method allows to have more stability of the gradient between batches and ease the training of the network. The batch size was set to 128, the learning rate starts at 0.001 with a learning rate scheduler that diminish the learning by tenfold if the metrics has not been improved in the last 10 epoch.

Additionally, dropout functions, functions that mask the output of previous layers were incorporated into our neural network to support it to avoid over-fitting.

Concerning the input of the model, the data was split into training, dev, and test set using with 80%, 10%, 10% of the overall data set. In the training phase, each instance of a batch we sampled randomly a companies of the training set, sampled a 3 months (63 days) of observation of the companies' stock returns along with the returns of the sector indices.

7 Benchmark

Intuitively, we could use random guessing as a benchmark (this would yield a 13% accuracy at best as Financial Services is the biggest represented class). A bit more challenging, we could use correlation as measure of association and using the sector with the highest correlation to our input series, This classifier gets an accuracy rates of 59%.

8 Methodology

8.1 Data processing

From the Quandl dataset, the preprocessing involves keeping only the ticker and the close price for as many date as possible and as many companies as possible. Then the table is joined to the SimFin dataset containing sectors for 726 companies. In total we have 16 sectors, from which we can extract data. Due to the lack of GICS sectors, a few sectors were merged together to increase their size, e.g. all the Oil and Gas companies were merged into a single sector.

Then the sector indices were created by averaging the daily returns of the stocks within the sectors. The returns were floor and capped to 10% as it is unlikely that a indices of stocks lose or gain more than 5% in a single trading day and the 100 level was set for 1990-01-01.

8.2 Implementation

The integrity of the code follows a linear process in the `notebook` folder of the project. One should be able to run all the code in each notebook separately. It was considered to have a proper implementation in the project and to avoid code duplication, but under time constraint, copy paste solution were preferred. That being said, the code has been written using pure python functions to avoid spaghetti code.

The implementation using Tensorflow and the keras API linked in the library. The Keras API allowed to defined our network and our model into a simple function and wrap customized transformation into the `keras.Lambda` layer. The exact implementation can be found in the notebooks. We launched AWS server with spot instance and launched a jupyter server there and made it accessible to our web-browser. We develop code also in the terminal with Emacs to adapt some code.

The first step was to download the data from the several providers and to process them as discussed in the previous section. Namely, we need to construct `pandas` dataframe (modeled as data frame with time index) of returns of stocks and synchronize theses returns with the

returns of the sectors. We decided to simply create a single big data frame with repeated values as the size of our data could largely be hold in memory.

The model is fed with 63 days of observations of stocks returns with sector returns in order predict one of 17 classes. The implementation of the model is defined through `keras` layers and custom functions. One particular custom function is to multiply the standardized returns of the stock with the sectors returns date by date to create the non-linear features. Keras function (models) only accept in a tensor/array format, hence the previous data frame need to be transformed into numpy array. In practice, we engineer three manual features to the model, the multiplication of the returns date by date, the sign of the previous operation and independent monthly correlation. These engineered features are then feed to different fully connected layers with relu activation, their results are concatenated and fed to a dense layer which creates the embedding, from which a last layer will create a partition for the classification.

In order to speed up the training phase, we needed to use several classes from keras to support asynchronous loading of the data thanks to the `Sequence` object. During training, batches were creating by sampling 128 sectors uniformly with resampling and select a random company (with uniform probability) in the selected sectors, this insure the classes are balanced as some sectors contains more stocks and ease the training.

As a technical detail, whenever the analysis and the training of the model occur in different scripts, the label encoder from scikit which maps the sectors into number should be at best serialized or insure its inputs are similar between the running time. Otherwise, there is a significant risk that the label encoder swap classes and destroy the perform of the classifier.

9 Refinement

The implementation has been performed with a simple function defining the network. We ran several experiment of the network, using convolutional networks, with adaptation of inception units and residual units, known in the neural network for images, but they did not lead to any improvement of the model. Moreover, to avoid overfitting, we added several batch normalization layers as well as Gaussian noise layer with a really small standard deviation. A few layers in the network were penalized L_2 regularization to insure that the features stayed as independent as possible.

The best models were the ones which were fed with correlations and forwarded to into a dense networks. The reason is associative measure a non linear and it is not common to multiply inputs with each other. That being said, we managed to create a convolutional layer that achieved an accuracy rate of 55%, a bit short from our best model and from the benchmark, but using only linear transformation.

9.1 Structure of the network

The network is depicted in Figure 9.1. From the input data, three transformation are performed. The first one create product of normalized observation in order to let the model to detect smaller pattern of interaction. The second transformation performs the same computation but on the sign of the input. This should create a more robust estimate of measure of association. The third is to compute the correlation matrix as feature for the model. We concatenate them and create a dense layer for creating the embedding from which we extract the classes.

10 Results

The base model using only correlation for the period of 3 months achieves 59*reallygood*.

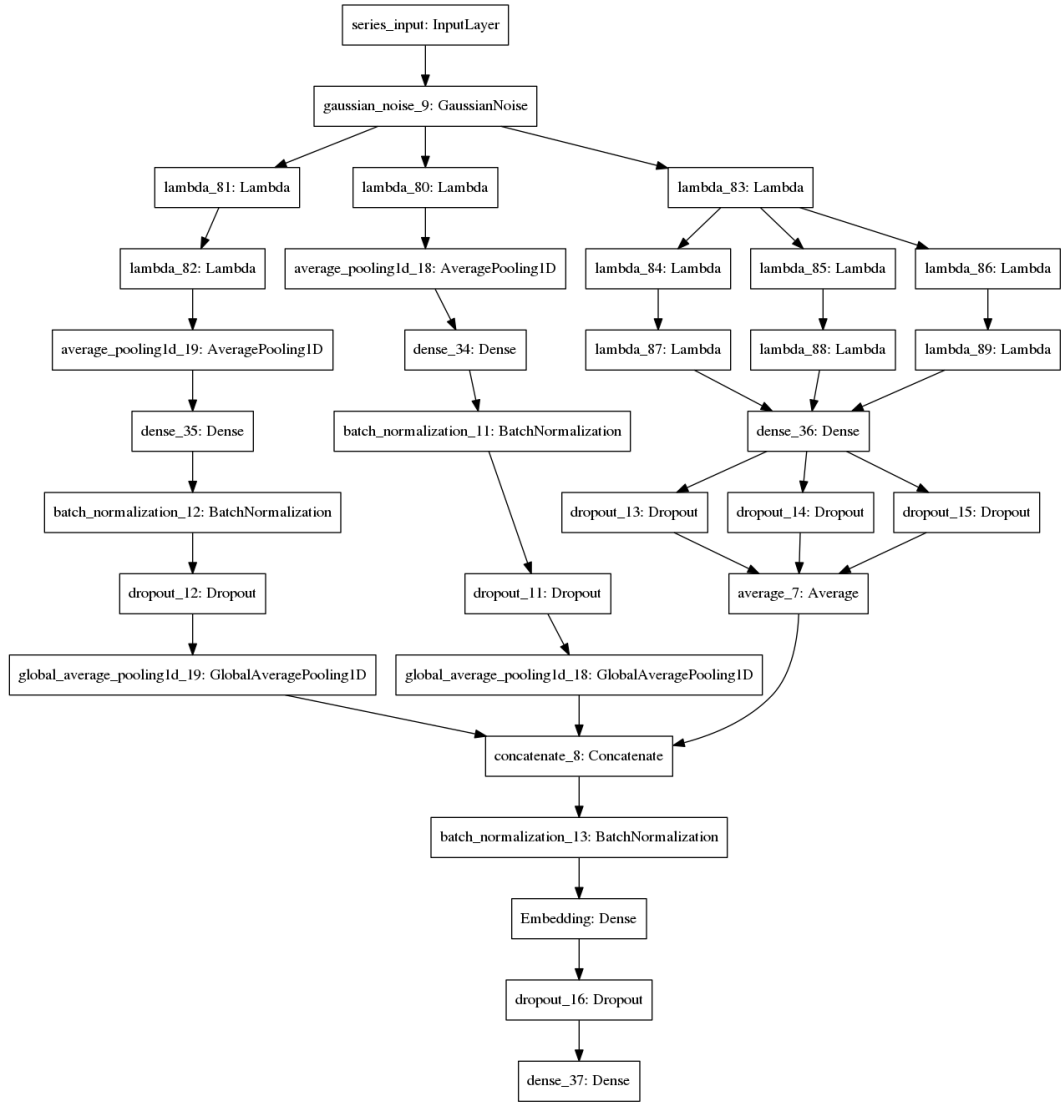


Figure 3: Neural Network structure

As for neural network model, it achieves around 58% percent accuracy on a single observation of three months which is on par with our benchmark. However, when we provide 25 random samples of 3 months period to the classifier, the classifier achieves 80% accuracy. As it can be read in Table 1




Figure 4: Confusion matrix of our predictor.

In Figure 10, we observe that the neural network model classifier does a fairly good job at classifying sectors with a notable exception of *Chemicals* and *Manufacturing - Apparels and Furniture*. The reason are probably that are little data.

The model is robust to input data as the input is standardized in the first step of the model and even a non-parametric feature (which is invariant to scale) is used in our model. As our training method heavily relies on sampling time period, stocks and dropout nodes, the optimal model is stable with respect to the random seed.

The model's performance generalize well to unseen stocks as the numbers in the in Table 1 and 2 corroborate. These Tables displays the performance of the classifier to unseen companies over different periods of time. Table 2 shows the performance of the classifier when it predicts the class of a company by aggregating the prediction over 25 sample periods.

Table 1: Confusion Report from the neural network classifier with 21 days observation data of companies from the test set.

	precision	recall	f1-score	support
Business Services	0.74	0.26	0.39	4474
Chemicals	0.63	0.55	0.59	2662
Communication Equipment	0.41	0.52	0.46	926
Communication Services	0.52	0.44	0.47	2307
Consumer Packaged Goods	0.59	0.51	0.54	3715
Drug Manufacturers	0.47	0.31	0.37	1603
Entertainment	0.56	0.56	0.56	2344
Financial Services	0.68	0.69	0.69	9241
Industrial Products	0.46	0.49	0.48	4091
Insurance	0.59	0.34	0.43	2807
Manufacturing - Apparel & Furniture	0.80	0.49	0.61	2031
Medical	0.48	0.67	0.56	8339
Oil and Gas	0.80	0.62	0.70	8897
REITs	0.75	0.69	0.72	5815
Retail - Apparel & Specialty	0.42	0.63	0.50	5740
Technology	0.57	0.60	0.59	8122
Utilities	0.64	0.95	0.77	3686
avg / total	0.61	0.59	0.59	76800

Table 2: Confusion Report from the neural network classifier with resampled data of 21 days from companies from the test set.

	precision	recall	f1-score	support
Business Services	1.00	0.67	0.80	3
Chemicals	1.00	1.00	1.00	3
Communication Equipment	1.00	1.00	1.00	2
Communication Services	1.00	0.50	0.67	2
Consumer Packaged Goods	1.00	1.00	1.00	4
Drug Manufacturers	0.00	0.00	0.00	2
Entertainment	1.00	1.00	1.00	3
Financial Services	1.00	1.00	1.00	10
Industrial Products	1.00	1.00	1.00	5
Insurance	1.00	1.00	1.00	3
Manufacturing - Apparel & Furniture	1.00	1.00	1.00	3
Medical	0.67	1.00	0.80	6
Oil and Gas	1.00	1.00	1.00	7

Continued on next page

Continued from previous page

	precision	recall	f1-score	support
REITs	1.00	1.00	1.00	6
Retail - Apparel & Specialty	1.00	0.83	0.91	6
Technology	0.80	1.00	0.89	8
Utilities	1.00	1.00	1.00	4
avg / total	0.93	0.94	0.92	77

10.1 Benchmark comparison

As mentioned previously, the model with no opinion (or constant answer) has an accuracy of at best 12% (the proportion of financial companies). Using the rule based method of choosing the class with the highest correlation, the model would yield an accuracy of around 59%. In light of these number, our classifier for a single period is slightly better (2% better accuracy) versus the rule based method, but it has the advantage of yielding embedding and also incorporate several measure of dependency. Moreover, it has the possibility to incorporate additional dependence measure that just the highest correlation.

11 Embedding

We are curious to look at the embedding produce by our neural network. We use t-SEN to create a low dimensional representation of it. This technique preserves the similarity between points.

In order to create an estimate of the embedding, we sampled the 25 periods of 3 months of each stock and averaged their embedding.

In Figure 11, it can be observed that stocks from the same sector tends to be near each other. The distinct cluster are finance and technology, which are also the most represented in our dataset. The model seems to have difficulty to differentiate some chemical companies as their embedding seems to be close to some industrial production companies. In general, the more companies we had in the raw dataset the more precise the groups are.



Figure 5: T-SNE of the embedding layers of the network.

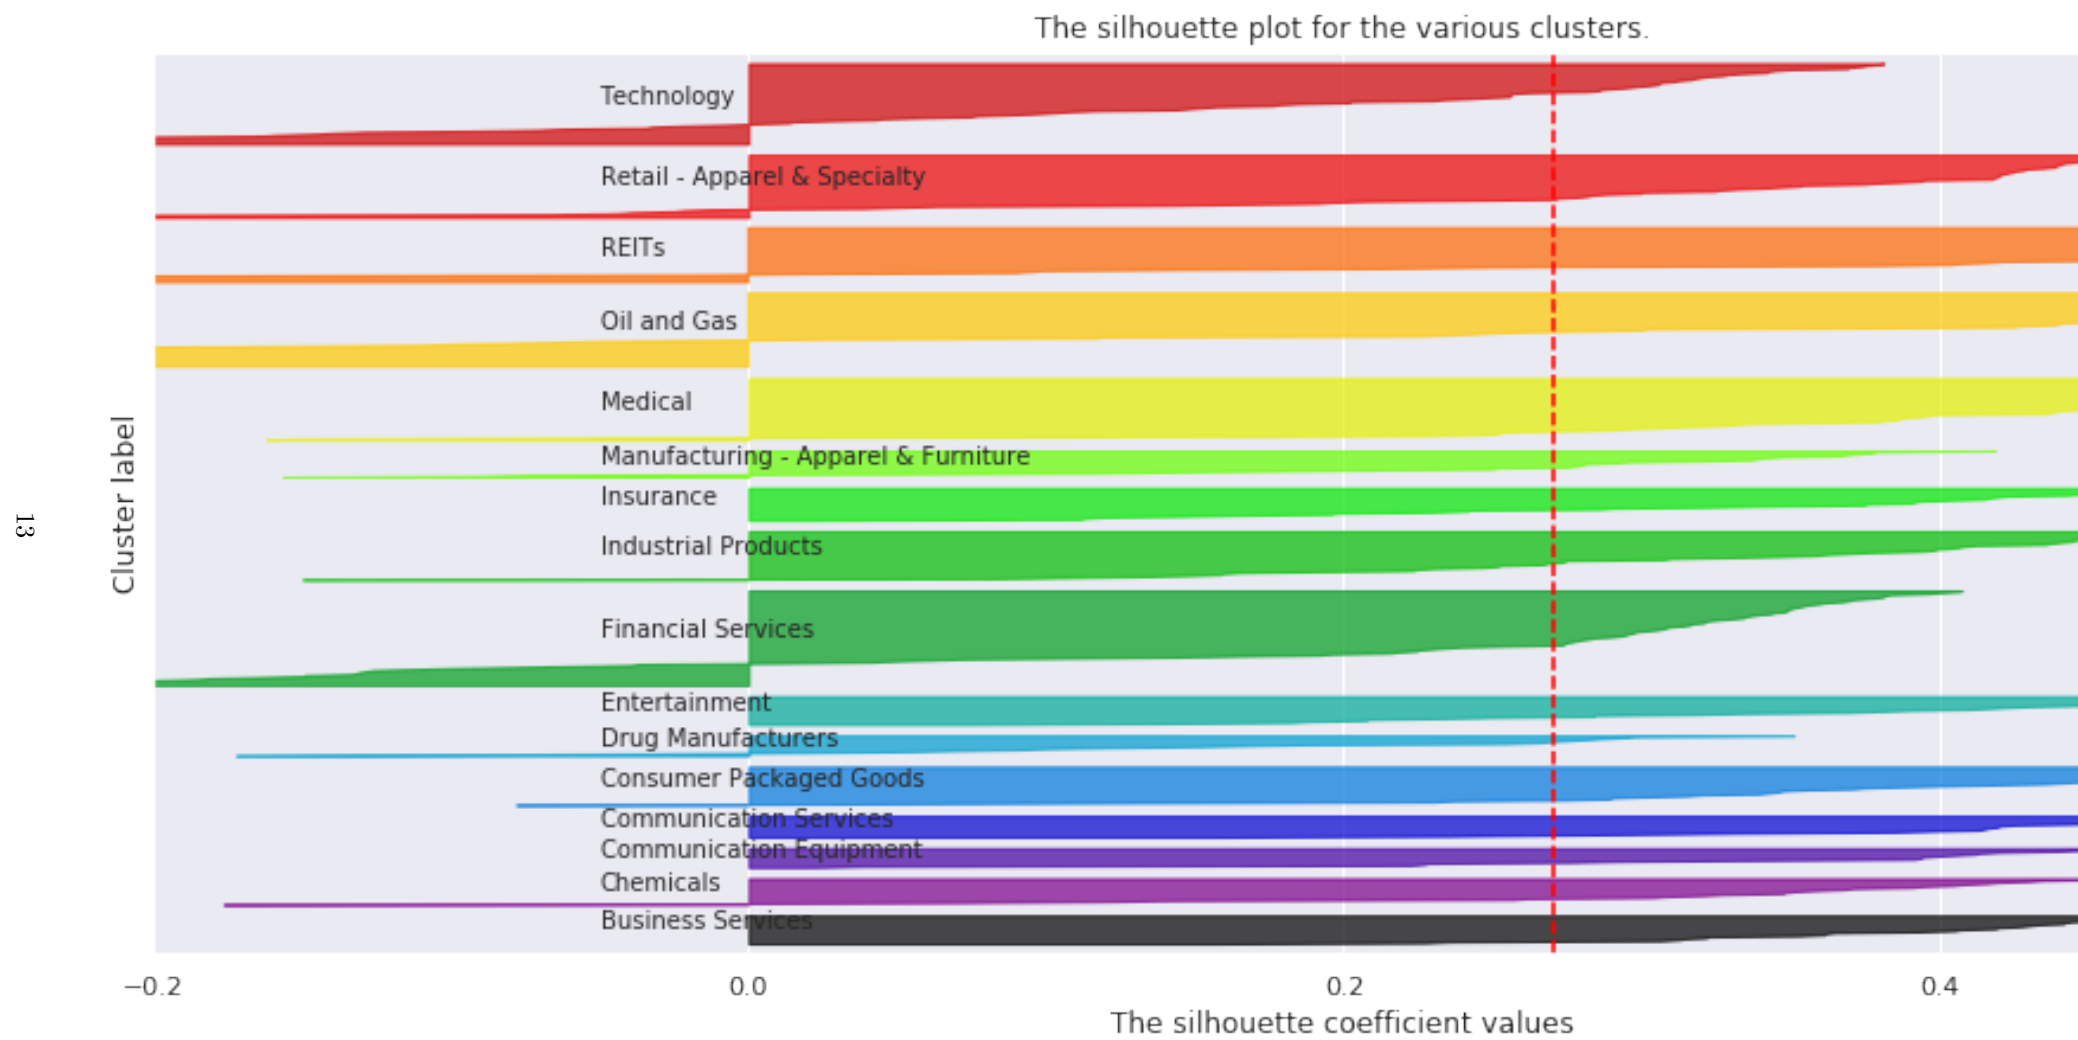


Figure 6: Silhouette score of the average embedding of the stocks

In Figure 11, the silhouette score is depicted for the several sectors in our dataset. The first observation is we see the advantage of a neural network classifier over k-means clustering, because the silhouette score over-weights the mislabeled sample of financial and technology, which are the most precise sectors. Otherwise we can observe that REITS, Oil and Gas, and the Insurance sector are grouped tightly making them quite distinct group.

12 Conclusion and Further applications

The first lesson I learned is data preparation and acquisition is much harder than thought and we should thank the machine learning community for providing so many labeled data set for our development. Indeed the financial industry still leverage on providing exclusive data and create a difficult task to leverage on alternative dataset, which could potentially provide added value.

Second, in deep learning, a bigger network does not necessarily translate into a better performance: training is much more difficult with more parameters even with regularizers and advanced optimization method. As for the training, balancing the classes improves a lot the training and can potentially improve the performance of the model. Financial data also proved to be tricky to handle without proper averaging. The signal to noise ratio is much higher than typical machine learning application domain.

The primary goal of the project was to find method that could create embedding for financial time series and Figure ?? provides a good proof that this goal has been reached. Our clustering abilities are still lacking as the accuracy rate for sample of three months is yet not better than a simple correlation measure. But the method has a higher accuracy when sampled with more data.

As for the improvement, we could try different architecture (LSTM and several skip convolution). The LSTM models could allow flexible time input. Moreover it would have been interesting to apply some semi-supervised method to improve the model and embedding. We could have applied our existing predictor for stocks whose sector is missing and recompute the indices and maybe retrain the classifier. A really interesting step would have been to incorporate T-SNE or Uniform Manifold Approximation and Projection in the training and apply additional clustering technique. These low dimensional projections seems to cluster data efficiently. This could also potentially resolve our inability to detect new group in as we would need to train them. Zero shot learning would be an interesting project to the study.